

API Rate Limits & Quotas

Per-Tier Rate Limits

Starter tier: 100 requests per minute, 5,000 per day
Growth tier: 500 requests per minute, 50,000 per day
Enterprise tier: 2,000 requests per minute, unlimited per day
All tiers share a burst allowance of 2x the per-minute limit for up to 10 seconds before throttling applies.

Rate Limit Headers

Every API response includes:
X-RateLimit-Limit: your per-minute ceiling
X-RateLimit-Remaining: requests left this window
X-RateLimit-Reset: Unix timestamp when the window resets
When you hit the limit the response is HTTP 429 with a
Retry-After header giving seconds until the window resets.

Concurrency Limits

Starter: 5 concurrent requests
Growth: 20 concurrent requests
Enterprise: 100 concurrent requests (configurable)
Concurrent requests beyond the limit receive HTTP 429 with a Retry-After of 1 second.

Token Quotas (LLM endpoints)

Starter: 500,000 input tokens / month
Growth: 5,000,000 input tokens / month
Enterprise: Negotiated -- contact sales@example.com
Output tokens are not counted against the quota.
Quota resets on the 1st of each calendar month UTC.
Overage: requests above quota return HTTP 402 with a link to the upgrade page.

Exemptions

Webhook delivery callbacks are not subject to rate limits.
Internal service-to-service calls using machine API keys (role: service_account) are exempt from per-minute limits but are capped at the daily maximum.